

Subset-Sum Hash Specification

Yossi Gilad, David Lazar, and Chris Peikert

Algorand, Inc.

September 3, 2021

1 Notation and Background

All logarithms are base two, i.e., $\log = \log_2$, unless otherwise indicated by a different subscript.

Binary strings. For a positive integer L , $\{0, 1\}^{<L}$ denotes the set of binary strings of length strictly less than L (including the zero-length empty string ε), and $\{0, 1\}^*$ denotes the set of all binary strings of any length. For binary strings u, v of any length, uv denotes their concatenation. For a binary string $x \in \{0, 1\}^*$, $|x| \geq 0$ denotes its length in bits. For a positive integer e and a non-negative integer $z < 2^e$, $\langle z \rangle_e \in \{0, 1\}^e$ denotes the little-endian representation of z as a binary string of exactly e bits.

Modular integers. For a positive integer q , $\mathbb{Z}_q := \mathbb{Z}/q\mathbb{Z}$ denotes the (abelian) ring of integers modulo q . Formally, \mathbb{Z}_q is the set of *cosets* of the form

$$c = \tilde{c} + q\mathbb{Z} := \{\dots, \tilde{c} - 2q, \tilde{c} - q, \tilde{c}, \tilde{c} + q, \tilde{c} + 2q, \dots\}$$

for some integer $\tilde{c} \in \mathbb{Z}$, with $\tilde{c} + q\mathbb{Z} = \tilde{c}' + q\mathbb{Z}$ if and only if q divides $\tilde{c} - \tilde{c}'$. In general, a coset can be represented by any of its elements, which might even differ from one location to the next. To avoid any ambiguity, this specification always requires a coset $c \in \mathbb{Z}_q$ to be externally read and written using its (unique) *distinguished representative* $\bar{c} \in \{0, 1, \dots, q-1\} \cap c$.

Vectors and matrices. Vectors are denoted by lower-case bold letters (e.g., \mathbf{a}), and matrices by upper-case bold letters (e.g., \mathbf{A}). *In this specification, vector and matrix entries are always indexed starting from zero.* A vector's i th entry is denoted by the same lower-case letter, but without boldface, and with a subscript i ; e.g., x_i is the i th entry of \mathbf{x} . Similarly, the i th column of a matrix is denoted by the same letter, but in lower case, and with a subscript i ; e.g., \mathbf{a}_i is the i th column of \mathbf{A} .

The set of n -dimensional vectors over a set X is denoted X^n . In particular, \mathbb{Z}_q^n is an (abelian) additive group, where the group operation is coordinate-wise addition (modulo q). Similarly, $X^{n \times m}$ denotes the set of n -by- m -dimensional matrices over X . For convenience, this specification sometimes uses standard vector and matrix operations (like sums and products) where they are well defined.

2 Compression Function Family

This section gives the mathematical definition of the subset-sum compression function family, which was first studied and popularized in early works like [IN89, Ajt96]. For specific parameters, the concrete security of the function against various kinds of (classical and quantum) cryptanalytic attacks is analyzed in a separate work.

2.1 Parameters

The subset-sum compression function family is parameterized by:

- a positive integer *modulus* $q \in \mathbb{N}$ (often taken to be a power of two);
- a positive integer *dimension* $n \in \mathbb{N}$;
- a positive integer *input length* $m \in \mathbb{N}$, where $m > n \log q$.

2.2 Function Definition

The compression function family for parameters q, n, m is defined as the collection

$$\mathcal{F}_{q,n,m} := \{f_{\mathbf{A}} : \{0, 1\}^m \rightarrow \mathbb{Z}_q^n : \mathbf{A} \in \mathbb{Z}_q^{n \times m}\},$$

where each function $f_{\mathbf{A}} \in \mathcal{F}_{q,n,m}$ is defined as

$$f_{\mathbf{A}}(\mathbf{x}) := \mathbf{A}\mathbf{x} = \sum_{i=1}^m x_i \cdot \mathbf{a}_i = \sum_{i:x_i=1} \mathbf{a}_i \in \mathbb{Z}_q^n. \quad (2.1)$$

The latter summation explains the name “subset-sum hash”: the output is the subset-sum of the columns of \mathbf{A} indicated by the bits of the input \mathbf{x} .

Observe that the condition $m > n \log q$ ensures that the functions in the family are *compressing*, i.e., the cardinality of their common domain $\{0, 1\}^m$ is strictly larger than that of their common range \mathbb{Z}_q^n : $2^m > 2^{n \log q} = q^n$.

When $q = 2^u$ is a power of two, $\ell := n \log q = nu$ is called the *output length*, and $b := m - \ell > 0$ is called the *block length*. In this case, for any $\mathbf{y} \in \mathbb{Z}_q^n$, let $\langle \mathbf{y} \rangle \in \{0, 1\}^\ell$ denote the representation of \mathbf{y} as an ℓ -bit string, obtained as the concatenation of the (u -bit, little-endian representations of the) distinguished representatives $\bar{y}_i \in \{0, 1, \dots, q-1\}$ of the coordinates y_i :

$$\langle \mathbf{y} \rangle := \langle \bar{y}_0 \rangle_u \langle \bar{y}_1 \rangle_u \cdots \langle \bar{y}_{n-1} \rangle_u \in \{0, 1\}^\ell. \quad (2.2)$$

2.3 Security Properties

Conjectured properties. For appropriate parameters, the subset-sum compression function family $\mathcal{F}_{q,n,m}$ is conjectured to have the following security properties:

- *Uninvertibility (UI)*: given uniformly random and independent $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ and $\mathbf{y} \in \mathbb{Z}_q^n$, it is infeasible to find some $\mathbf{x} \in \{0, 1\}^m$ such that $f_{\mathbf{A}}(\mathbf{x}) = \mathbf{A}\mathbf{x} = \mathbf{y}$.
- *One-wayness (OW)*: given $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ and $\mathbf{y} = f_{\mathbf{A}}(\mathbf{x}) = \mathbf{A}\mathbf{x} \in \mathbb{Z}_q^n$ (but not \mathbf{x} itself), where $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ and $\mathbf{x} \in \{0, 1\}^m$ are uniformly random and independent, it is infeasible to find some $\mathbf{x}' \in \{0, 1\}^m$ (not necessarily different from \mathbf{x}) such that $f_{\mathbf{A}}(\mathbf{x}') = \mathbf{y}$.

- *Second-preimage resistance (SPR)*: given uniformly random and independent $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ and $\mathbf{x} \in \{0, 1\}^m$, it is infeasible to find some $\mathbf{x}' \in \{0, 1\}^m$ such that $\mathbf{x}' \neq \mathbf{x}$ and $f_{\mathbf{A}}(\mathbf{x}') = f_{\mathbf{A}}(\mathbf{x})$.
- *Target-collision resistance (TCR)*: it is infeasible to choose some $\mathbf{x} \in \{0, 1\}^m$ and then, given a uniformly random $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$, to find some distinct $\mathbf{x}' \in \{0, 1\}^m \setminus \{\mathbf{x}\}$ such that $f_{\mathbf{A}}(\mathbf{x}) = f_{\mathbf{A}}(\mathbf{x}')$, i.e., $\mathbf{Ax} = \mathbf{Ax}' \in \mathbb{Z}_q^n$.
- *Collision resistance (CR)*: given a uniformly random $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$, it is infeasible to find distinct $\mathbf{x}, \mathbf{x}' \in \{0, 1\}^m$ such that $f_{\mathbf{A}}(\mathbf{x}) = f_{\mathbf{A}}(\mathbf{x}')$, i.e., $\mathbf{Ax} = \mathbf{Ax}' \in \mathbb{Z}_q^n$.

Note that breaking CR is equivalent to finding a nonzero $\mathbf{z} \in \{-1, 0, 1\}^m \setminus \{\mathbf{0}\}$ such that $\mathbf{Az} = \mathbf{0} \in \mathbb{Z}_q^n$. In one direction, given such \mathbf{x}, \mathbf{x}' , define $\mathbf{z} = \mathbf{x} - \mathbf{x}' \in \{-1, 0, 1\}^m \setminus \{\mathbf{0}\}$ and observe that $\mathbf{Az} = \mathbf{Ax} - \mathbf{Ax}' = \mathbf{0}$. In the other direction, given such \mathbf{z} , let $\mathbf{x} \in \{0, 1\}^m$ be 1 (respectively, 0) wherever \mathbf{z} is 1 (resp., 0 or -1), and similarly let $\mathbf{x}' \in \{0, 1\}^m$ be 1 (respectively, 0) wherever \mathbf{z} is -1 (resp., 0 or 1). Then $\mathbf{z} = \mathbf{x} - \mathbf{x}'$, and since $\mathbf{0} = \mathbf{Az} = \mathbf{A}(\mathbf{x} - \mathbf{x}')$, we have $\mathbf{Ax} = \mathbf{Ax}'$, as desired.

It is well known that CR tightly implies TCR generically (i.e., for any compression function family), because breaking TCR immediately yields a collision. Moreover, TCR tightly implies SPR generically, because being able to find a collision with a uniformly random input (that is independent of \mathbf{A}) in particular means being able to find a collision with an input of one's choice (since that input may just be chosen uniformly). Furthermore, for parameters that yield significant compression (i.e., for $m \gg n \log q$), it is well known that SPR tightly implies OW generically, and that OW implies UI for the subset-sum family.

While none of the converse directions is known to hold generically, for the subset-sum family with significant compression, a *relaxed* form of UI implies CR, thus implying that all four security properties are closely related. Specifically, any attack that breaks CR with probability δ can be used to successfully attack relaxed-UI in essentially the same amount of time and with probability $\approx \delta/(2m)$, where in relaxed-UI the output \mathbf{x} may be taken from $\{-1, 0, 1\}^m$ (it is not limited to $\{0, 1\}^m$).

Non-conjectured properties. On the other hand, the family $\mathcal{F}_{q,n,m}$ is *not conjectured to have*, or is even *known not to have*, the following security properties:

- *Unpredictability/Pseudorandomness*: outputs of $f_{\mathbf{A}}$ on different inputs do not appear random or uncorrelated, even if parts of the corresponding inputs are unknown to the attacker. This is because $f_{\mathbf{A}}$ is linear:

$$f_{\mathbf{A}}(\mathbf{x} + \mathbf{x}') = \mathbf{A}(\mathbf{x} + \mathbf{x}') = \mathbf{Ax} + \mathbf{Ax}' = f_{\mathbf{A}}(\mathbf{x}) + f_{\mathbf{A}}(\mathbf{x}'),$$

as long as $\mathbf{x}, \mathbf{x}', \mathbf{x} + \mathbf{x}' \in \{0, 1\}^m$, which is easy to arrange in many contexts. In particular, any 0 bits of \mathbf{x} can be changed to 1s by adding an \mathbf{x}' that is 1 in the suitable position(s), and 0 wherever \mathbf{x} is 1.

- *Random oracle*: for the same reasons as above, $f_{\mathbf{A}}$ does not “behave like a random oracle” in many of the ways that are typically expected. Therefore, it should not be used to instantiate a random oracle in any cryptosystems or protocols.

3 Hashing Arbitrary-Length Messages

The compression functions defined in Section 2 map a *fixed-length* m -bit input to an output of length $\ell := n \log q < m$. As is standard, a message of *arbitrary* length is hashed to a fixed-length output by invoking the compression function one or more times, using the Merkle–Damgård (MD) transform [Mer89, Dam89].

3.1 Padding

The transform uses the following padding method, which maps a binary string of any bounded length into a strictly longer one whose length is a multiple of the block length $b := m - \ell > 0$. Formally, for any positive integer e , define the padding function $\text{pad}_{b,e} : \{0, 1\}^{<2^e} \rightarrow \{0, 1\}^*$ as

$$\text{pad}_{b,e}(x) = x10^z \langle |x| \rangle_e, \quad (3.1)$$

where $z \geq 0$ is the smallest non-negative integer for which $|x| + 1 + z + e$ is a multiple of b . That is, $r := b \cdot \lceil k/b \rceil - k \in \{0, 1, \dots, b-1\}$, where $k = |x| + 1 + e$. Note especially that $\text{pad}_{b,e}(x)$ *unconditionally* appends at least $1 + e$ bits to x , even if $|x|$ itself is a multiple of b .

The purpose of the padding function is to produce a string whose length is a multiple of the block length, and so that the MD transform preserves collision resistance. That is, any collision in the full hash function immediately yields a collision in the underlying compression function.

3.2 Hash Functions

Fix subset-sum parameters q, n, m where $q = 2^u$ is a power of two for some positive integer u . Let $\ell := n \log q = nu$ be the output length in bits, and let $b := m - \ell > 0$ be the block length.

This specification defines two functions on arbitrary-length input strings: an “unsalted” mode that takes only the string as input, and a “salted” mode, originally defined in [HK06], that additionally takes a public salt value (which should typically be chosen uniformly at random, or pseudorandomly).

3.2.1 Unsalted Mode

For any matrix $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ defining the compression function $f_{\mathbf{A}} : \{0, 1\}^m \rightarrow \mathbb{Z}_q^n$ (Section 2) and a positive integer $e > 0$, define the unsalted hash function

$$\begin{aligned} H_{\mathbf{A},e} : \{0, 1\}^{<2^e} &\rightarrow \{0, 1\}^\ell \\ H_{\mathbf{A},e}(x) &:= H'_{\mathbf{A}}(0^\ell, \text{pad}_{b,e}(x)), \end{aligned} \quad (3.2)$$

where the chaining function $H'_{\mathbf{A}} : \{0, 1\}^\ell \times (\bigcup_{i=0}^\infty \{0, 1\}^{ib}) \rightarrow \{0, 1\}^\ell$ is defined as

$$H'_{\mathbf{A}}(h, w) := \begin{cases} h & \text{if } w = \varepsilon \\ H'_{\mathbf{A}}(\langle f_{\mathbf{A}}(hu) \rangle, v) & \text{where } w = uv \text{ for } u \in \{0, 1\}^b. \end{cases} \quad (3.3)$$

(Recall that representation $\langle \mathbf{y} \rangle \in \{0, 1\}^\ell$ for $\mathbf{y} \in \mathbb{Z}_q^n$ is defined in Section 2.2.)

3.2.2 Salted Mode

For any matrix $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ defining the compression function $f_{\mathbf{A}} : \{0, 1\}^m \rightarrow \mathbb{Z}_q^n$ (Section 2) and a positive integer $e > 0$, define the salted hash function

$$\begin{aligned} \tilde{H}_{\mathbf{A},e} : \{0, 1\}^{<2^e} \times \underbrace{\{0, 1\}^b}_{\text{salt}} &\rightarrow \{0, 1\}^\ell \\ \tilde{H}_{\mathbf{A},e}(x, r) &:= \tilde{H}'_{\mathbf{A}}(0^\ell, \text{pad}_{b,e}(0^b x), r) \end{aligned} \quad (3.4)$$

where the chaining function $\tilde{H}'_{\mathbf{A}} : \{0, 1\}^\ell \times (\bigcup_{i=0}^\infty \{0, 1\}^{ib}) \times \{0, 1\}^b \rightarrow \{0, 1\}^\ell$ is defined as

$$\tilde{H}'_{\mathbf{A}}(h, w, r) := \begin{cases} h & \text{if } w = \varepsilon \\ \tilde{H}'_{\mathbf{A}}(\langle f_{\mathbf{A}}(h\tilde{u}) \rangle, v, r) & \text{where } w = uv \text{ for } u \in \{0, 1\}^b \text{ and } \tilde{u} = u \oplus r. \end{cases} \quad (3.5)$$

The only differences between the salted-mode functions $\tilde{H}_{\mathbf{A},e}$, $\tilde{H}'_{\mathbf{A}}$ and their corresponding unsalted-mode functions $H_{\mathbf{A},e}$, $H'_{\mathbf{A}}$ are depicted above with wavy underlines, and are as follows:

1. Each function $\tilde{H}_{\mathbf{A},e}$, $\tilde{H}'_{\mathbf{A}}$ takes an additional salt input $r \in \{0, 1\}^b$, whose length is the block length.
2. The function $\tilde{H}_{\mathbf{A},e}(x, r)$ prepends an all-zeros block $0^b \in \{0, 1\}^b$ to the input x (before padding it using $\text{pad}_{b,e}$ and inputting the result to the appropriate chaining function, as in the unsalted case). Note that this prepended zeros block counts toward the total input length in the padding function.
3. The chaining function $\tilde{H}'_{\mathbf{A}}$ XORs each block $u \in \{0, 1\}^b$ of the padded input—including the initial, prepended all-zeros block—with the salt value $r \in \{0, 1\}^b$ (before inputting the result together with the previous chaining value h to the compression function, as in the unsalted case).

4 Concrete Instantiations

The fully specified hash functions are merely instantiations of the unsalted function $H_{\mathbf{A},e}$ (Equation (3.2)) and salted function $\tilde{H}_{\mathbf{A},e}$ (Equation (3.4)) for specific subset-sum parameters q, n, m , matrices $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$, and input-length representation bit lengths e .

4.1 Deriving \mathbf{A}

It is well known that for typical parameters, it is easy to generate a matrix $\mathbf{A} \in \mathbb{Z}_q^{n \times m}$ that is indistinguishable from uniformly random, together with some *known collisions* in the associated subset-sum compression function $f_{\mathbf{A}}$. This may allow the party that generated \mathbf{A} to violate the security of constructions that use this function. Therefore, it is very important to generate a random-looking \mathbf{A} in a “nothing up my sleeve” manner that is highly unlikely to admit any such backdoor.

Let $q = 2^u$, n, m be subset-sum parameters where u is a positive integer, and let $\text{XOF} : \{0, 1\}^* \rightarrow \{0, 1\}^\infty$ represent a suitable cryptographic *extendable-output function*, such as SHAKE-256. Then for any identifier $id \in \{0, 1\}^*$, define the matrix $\mathbf{A}_{\text{XOF},id} \in \mathbb{Z}_q^{n \times m}$ as:

$$\mathbf{A}_{\text{XOF},id} := \text{pack}_{u,n,m}(\text{XOF}(\langle u \rangle_{16} \langle n \rangle_{16} \langle m \rangle_{16} id)_{0,\dots,unm-1}) \in \mathbb{Z}_q^{n \times m}, \quad (4.1)$$

where $\text{pack}_{u,n,m} : \{0, 1\}^{unm} \rightarrow \mathbb{Z}_q^{n \times m}$ constructs its output matrix from its input string in row-major order, using u bits per entry. That is, for all $i = 0, \dots, n-1$ and $j = 0, \dots, m-1$, the distinguished representative of the (i, j) th entry $a_{i,j} \in \mathbb{Z}_q$ of $\mathbf{A} = \text{pack}_{u,n,m}(w)$ has binary representation

$$\langle \bar{a}_{i,j} \rangle_u = w_{u(im+j), \dots, u(im+j)+(u-1)}.$$

4.2 Concrete Parameters

The implemented functions are (unsalted) $H := H_{\mathbf{A},e}$ and (salted) $\tilde{H} := \tilde{H}_{\mathbf{A},e}$, where:

- the modulus $q = 2^{64}$, so $u = \log q = 64 = 2^6$;
- the dimension $n = 8 = 2^3$, so the output length is $\ell = n \log q = 512 = 2^9$;
- the input length $m = 1024 = 2^{10}$, so the block length is $b = m - \ell = 512 = 2^9$;
- the representation length (of the hash input length) is $e = 128 = 2^7$;
- the extendable-output function is XOF = SHAKE-256;
- the matrix is $\mathbf{A} := \mathbf{A}_{\text{XOF},id}$ where id is the ASCII representation of Algorand.

References

- [Ajt96] M. Ajtai. Generating hard instances of lattice problems. *Quaderni di Matematica*, 13:1–32, 2004. Preliminary version in STOC 1996.
- [Dam89] I. Damgård. A design principle for hash functions. In *CRYPTO*, pages 416–427. 1989.
- [HK06] S. Halevi and H. Krawczyk. Strengthening digital signatures via randomized hashing. In *CRYPTO*, pages 41–59. 2006.
- [IN89] R. Impagliazzo and M. Naor. Efficient cryptographic schemes provably as secure as subset sum. *J. Cryptology*, 9(4):199–216, 1996. Preliminary version in FOCS 1989.
- [Mer89] R. C. Merkle. A certified digital signature. In *CRYPTO*, pages 218–238. 1989.